

```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

import statsmodels.formula.api as smf
from statsmodels.graphics.regressionplots import influence_plot
```

```
In [2]: import warnings
warnings.filterwarnings('ignore')
```

```
In [3]: #Read the data
cars = pd.read_csv("Cars.csv")
cars.head()
```

```
Out[3]:
```

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149

## Correlation Matrix

```
In [4]: cars.corr()
```


```
Out[4]:
```

	HP	MPG	VOL	SP	WT
HP	1.000000	-0.725038	0.077459	0.973848	0.076513
MPG	-0.725038	1.000000	-0.529057	-0.687125	-0.526759
VOL	0.077459	-0.529057	1.000000	0.102170	0.999203
SP	0.973848	-0.687125	0.102170	1.000000	0.102439
WT	0.076513	-0.526759	0.999203	0.102439	1.000000

## Scatterplot between variables along with histograms

```
In [5]: #Format the plot background and scatter plots for all the variables
sns.set_style(style='darkgrid')
sns.pairplot(cars)
```

```
Out[5]: <seaborn.axisgrid.PairGrid at 0x7fe4193a8340>
```



## Preparing a model

```
In [6]: #Build model
import statsmodels.formula.api as smf
model = smf.ols('MPG~WT+VOL+SP+HP',data=cars).fit()
```

```
In [7]: model.summary()
```

Out[7]:

## OLS Regression Results

<b>Dep. Variable:</b>	MPG	<b>R-squared:</b>	0.771
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.758
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	63.80
<b>Date:</b>	Wed, 02 Feb 2022	<b>Prob (F-statistic):</b>	1.54e-23
<b>Time:</b>	20:29:08	<b>Log-Likelihood:</b>	-233.96
<b>No. Observations:</b>	81	<b>AIC:</b>	477.9
<b>Df Residuals:</b>	76	<b>BIC:</b>	489.9
<b>Df Model:</b>	4		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	30.6773	14.900	2.059	0.043	1.001	60.354
<b>WT</b>	0.4006	1.693	0.237	0.814	-2.972	3.773
<b>VOL</b>	-0.3361	0.569	-0.591	0.556	-1.469	0.796
<b>SP</b>	0.3956	0.158	2.500	0.015	0.080	0.711
<b>HP</b>	-0.2054	0.039	-5.239	0.000	-0.284	-0.127

<b>Omnibus:</b>	10.780	<b>Durbin-Watson:</b>	1.403
<b>Prob(Omnibus):</b>	0.005	<b>Jarque-Bera (JB):</b>	11.722
<b>Skew:</b>	0.707	<b>Prob(JB):</b>	0.00285
<b>Kurtosis:</b>	4.215	<b>Cond. No.</b>	6.09e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.09e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In [8]:

```
#Coefficients
model.params
```

Out[8]:

```
Intercept    30.677336
WT            0.400574
VOL          -0.336051
SP            0.395627
HP           -0.205444
dtype: float64
```

In [9]:

```
#t and p-Values
print(model.tvalues, '\n', model.pvalues)
```

```
Intercept    2.058841
WT           0.236541
VOL          -0.590970
SP           2.499880
HP           -5.238735
dtype: float64
Intercept    0.042936
WT           0.813649
VOL          0.556294
SP           0.014579
HP           0.000001
dtype: float64
```

```
In [10]: #R squared values
(model.rsquared,model.rsquared_adj)
```

```
Out[10]: (0.7705372737359842, 0.7584602881431413)
```

## Simple Linear Regression Models

```
In [11]: ml_v=smf.ols('MPG~VOL',data = cars).fit()
#t and p-Values
print(ml_v.tvalues, '\n', ml_v.pvalues)
ml_v.summary()
```

```
Intercept    14.106056
VOL          -5.541400
dtype: float64
Intercept    2.753815e-23
VOL          3.822819e-07
dtype: float64
```

Out[11]:

## OLS Regression Results

<b>Dep. Variable:</b>	MPG	<b>R-squared:</b>	0.280
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.271
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	30.71
<b>Date:</b>	Wed, 02 Feb 2022	<b>Prob (F-statistic):</b>	3.82e-07
<b>Time:</b>	20:29:08	<b>Log-Likelihood:</b>	-280.28
<b>No. Observations:</b>	81	<b>AIC:</b>	564.6
<b>Df Residuals:</b>	79	<b>BIC:</b>	569.4
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	55.8171	3.957	14.106	0.000	47.941	63.693
<b>VOL</b>	-0.2166	0.039	-5.541	0.000	-0.294	-0.139
<b>Omnibus:</b>	2.691	<b>Durbin-Watson:</b>	0.566			
<b>Prob(Omnibus):</b>	0.260	<b>Jarque-Bera (JB):</b>	1.997			
<b>Skew:</b>	-0.263	<b>Prob(JB):</b>	0.368			
<b>Kurtosis:</b>	3.562	<b>Cond. No.</b>	462.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [12]: ml_w=smf.ols('MPG~WT',data = cars).fit()
print(ml_w.tvalues, '\n', ml_w.pvalues)
ml_w.summary()
```

```
Intercept    14.248923
WT           -5.508067
dtype: float64
Intercept    1.550788e-23
WT           4.383467e-07
dtype: float64
```

Out[12]:

## OLS Regression Results

<b>Dep. Variable:</b>	MPG	<b>R-squared:</b>	0.277
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.268
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	30.34
<b>Date:</b>	Wed, 02 Feb 2022	<b>Prob (F-statistic):</b>	4.38e-07
<b>Time:</b>	20:29:08	<b>Log-Likelihood:</b>	-280.42
<b>No. Observations:</b>	81	<b>AIC:</b>	564.8
<b>Df Residuals:</b>	79	<b>BIC:</b>	569.6
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	55.2296	3.876	14.249	0.000	47.514	62.945
<b>WT</b>	-0.6420	0.117	-5.508	0.000	-0.874	-0.410
<b>Omnibus:</b>	2.735	<b>Durbin-Watson:</b>	0.555			
<b>Prob(Omnibus):</b>	0.255	<b>Jarque-Bera (JB):</b>	2.045			
<b>Skew:</b>	-0.263	<b>Prob(JB):</b>	0.360			
<b>Kurtosis:</b>	3.573	<b>Cond. No.</b>	149.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [13]: ml_wv=smf.ols('MPG~WT+VOL',data = cars).fit()
print(ml_wv.tvalues, '\n', ml_wv.pvalues)
ml_wv.summary()
```

```
Intercept    12.545736
WT            0.489876
VOL          -0.709604
dtype: float64
Intercept    2.141975e-20
WT           6.255966e-01
VOL          4.800657e-01
dtype: float64
```

Out[13]:

## OLS Regression Results

<b>Dep. Variable:</b>	MPG	<b>R-squared:</b>	0.282
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.264
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	15.33
<b>Date:</b>	Wed, 02 Feb 2022	<b>Prob (F-statistic):</b>	2.43e-06
<b>Time:</b>	20:29:08	<b>Log-Likelihood:</b>	-280.16
<b>No. Observations:</b>	81	<b>AIC:</b>	566.3
<b>Df Residuals:</b>	78	<b>BIC:</b>	573.5
<b>Df Model:</b>	2		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	56.8847	4.534	12.546	0.000	47.858	65.912
<b>WT</b>	1.4349	2.929	0.490	0.626	-4.397	7.266
<b>VOL</b>	-0.6983	0.984	-0.710	0.480	-2.658	1.261

<b>Omnibus:</b>	2.405	<b>Durbin-Watson:</b>	0.591
<b>Prob(Omnibus):</b>	0.300	<b>Jarque-Bera (JB):</b>	1.712
<b>Skew:</b>	-0.251	<b>Prob(JB):</b>	0.425
<b>Kurtosis:</b>	3.506	<b>Cond. No.</b>	597.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## Calculating VIF

```
In [14]: rsq_hp = smf.ols('HP~WT+VOL+SP',data=cars).fit().rsquared
vif_hp = 1/(1-rsq_hp) # 16.33

rsq_wt = smf.ols('WT~HP+VOL+SP',data=cars).fit().rsquared
vif_wt = 1/(1-rsq_wt) # 564.98

rsq_vol = smf.ols('VOL~WT+SP+HP',data=cars).fit().rsquared
vif_vol = 1/(1-rsq_vol) # 564.84

rsq_sp = smf.ols('SP~WT+VOL+HP',data=cars).fit().rsquared
vif_sp = 1/(1-rsq_sp) # 16.35

# Storing vif values in a data frame
d1 = {'Variables':['Hp','WT','VOL','SP'],'VIF':[vif_hp,vif_wt,vif_vol,vif_sp]}
Vif_frame = pd.DataFrame(d1)
Vif_frame
```

Out[14]:

	Variables	VIF
0	Hp	19.926589
1	WT	639.533818
2	VOL	638.806084
3	SP	20.007639

## Residual Analysis

### Test for Normality of Residuals (Q-Q Plot)

In [15]: `model.resid.min()`

Out[15]: -8.631961053868388

In [61]: `model.resid`


Out[61]:

```

0      10.258747
1       7.624608
2       7.734060
3       3.157963
4       8.331584
...
76     15.617904
77      1.298838
78      7.863547
79      7.517122
80     -3.458218
Length: 81, dtype: float64

```

In [16]: `import statsmodels.api as sm`  
`qqplot=sm.qqplot(model.resid,line='q') # line = 45 to draw the diagonal line`  
`plt.title("Normal Q-Q plot of residuals")`  
`plt.show()`



In [17]: `list(np.where(model.resid>10))`

Out[17]: `[array([ 0, 76])]`



# Residual Plot for Homoscedasticity

```
In [18]: def get_standardized_values( vals ):
          return (vals - vals.mean())/vals.std()
```

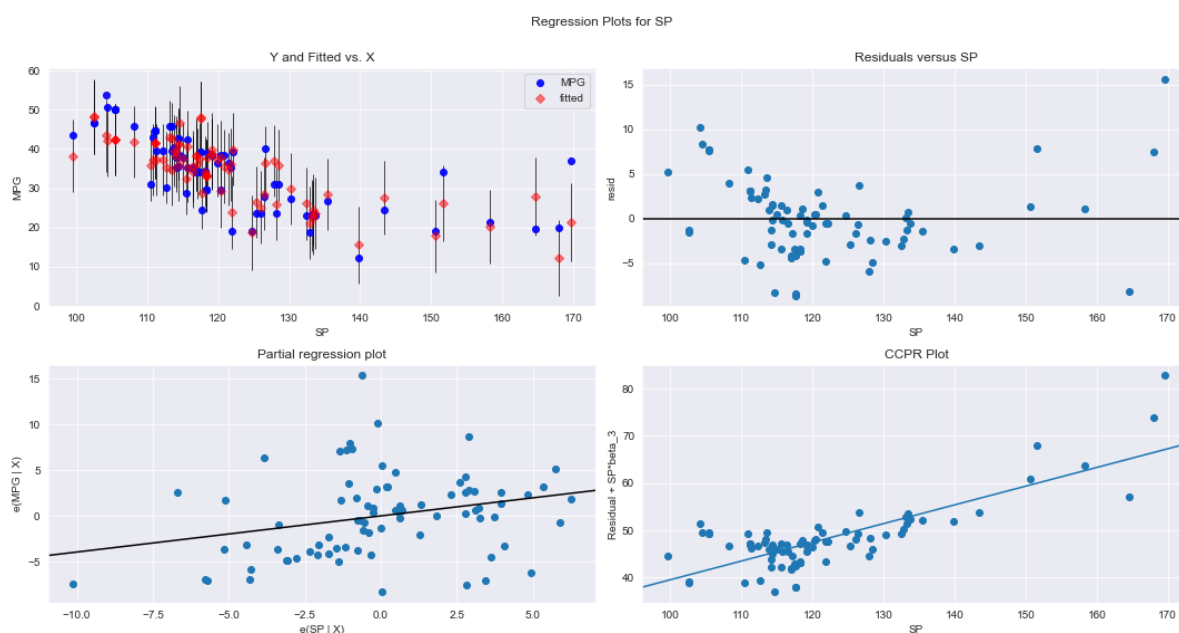
```
In [19]: plt.scatter(get_standardized_values(model.fittedvalues),
                    get_standardized_values(model.resid))

plt.title('Residual Plot')
plt.xlabel('Standardized Fitted values')
plt.ylabel('Standardized residual values')
plt.show()
```



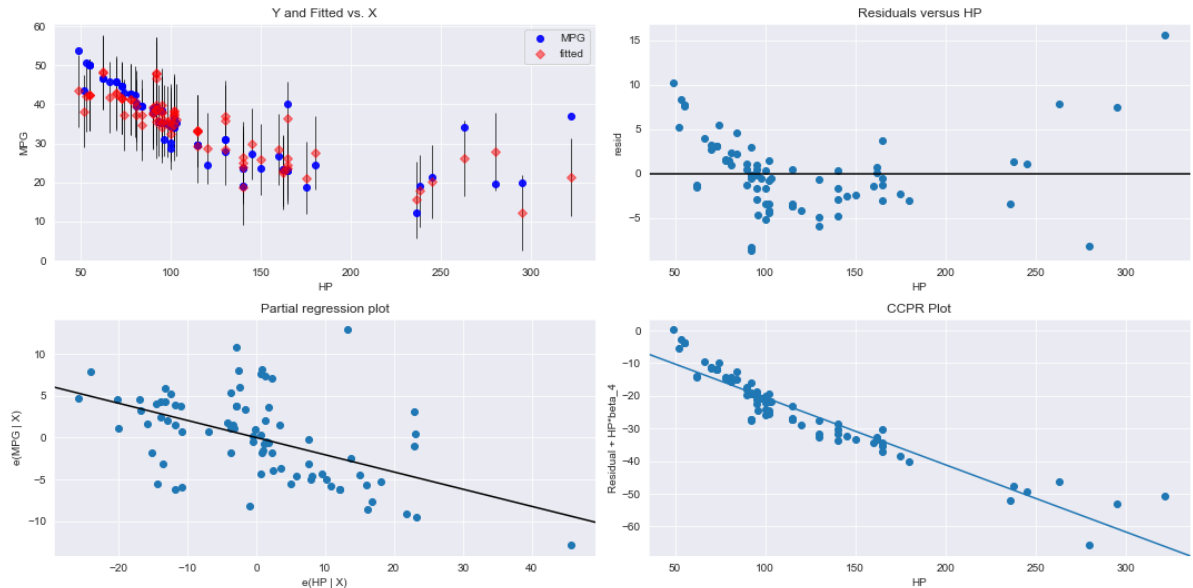
## Residual Vs Regressors

```
In [20]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model, "SP", fig=fig)
plt.show()
```



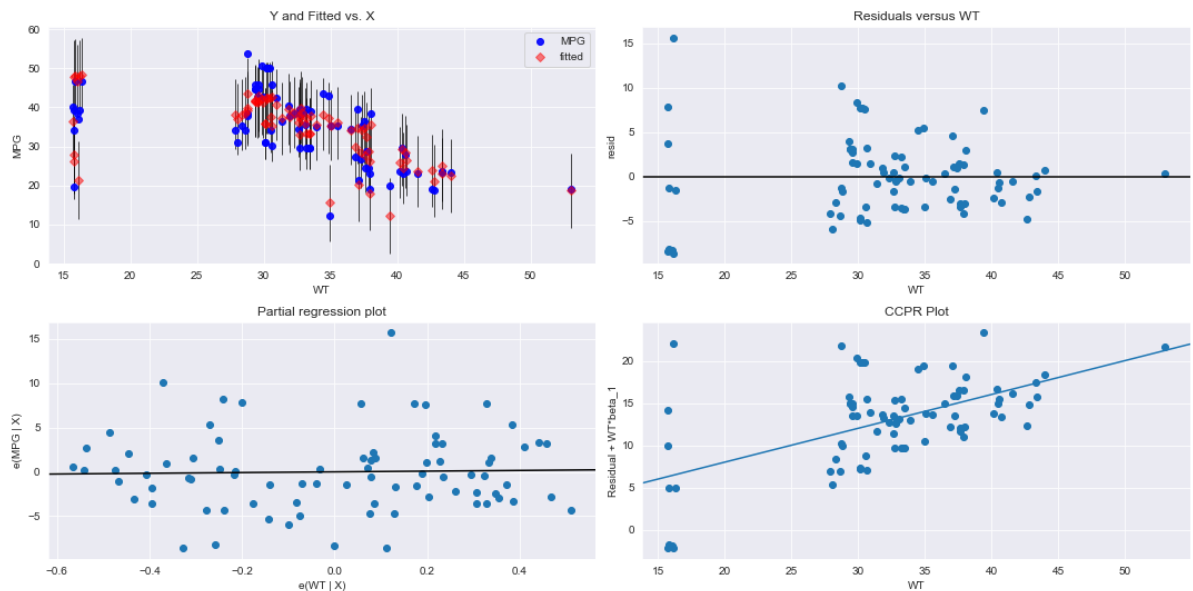
```
In [21]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model, "HP", fig=fig)
plt.show()
```

Regression Plots for HP



```
In [22]: fig = plt.figure(figsize=(15,8))
fig = sm.graphics.plot_regress_exog(model, "WT", fig=fig)
plt.show()
```

Regression Plots for WT



## Model Deletion Diagnostics


## Detecting Influencers/Outliers

## Cook's Distance

```
In [23]: model_influence = model.get_influence()
(c, _) = model_influence.cooks_distance
```

```
In [24]: #Plot the influencers values using stem plot
fig = plt.subplots(figsize=(20, 7))
plt.stem(np.arange(len(cars)), np.round(c, 3))
plt.xlabel('Row index')
```

```
plt.ylabel('Cooks Distance')
plt.show()
```




```
In [25]: #index and value of influencer where c is more than .5
         (np.argmax(c), np.max(c))
```

```
Out[25]: (76, 1.0865193998180052)
```

## High Influence points

```
In [26]: from statsmodels.graphics.regressionplots import influence_plot
         influence_plot(model)
         plt.show()
```



```
In [27]: cars.shape
```

```
Out[27]: (81, 5)
```

```
In [28]: k = cars.shape[1]
         n = cars.shape[0]
         leverage_cutoff = 3*((k + 1)/n)
         leverage_cutoff
```

```
Out[28]: 0.2222222222222222
```

From the above plot, it is evident that data point 70 and 76 are the influencers

```
In [29]: cars[cars.index.isin([70, 76])]
```

```
Out[29]:
```

	HP	MPG	VOL	SP	WT
<b>70</b>	280	19.678507	50	164.598513	15.823060
<b>76</b>	322	36.900000	50	169.598513	16.132947

```
In [30]: #See the differences in HP and other variable values  
cars.head()
```

```
Out[30]:
```

	HP	MPG	VOL	SP	WT
<b>0</b>	49	53.700681	89	104.185353	28.762059
<b>1</b>	55	50.013401	92	105.461264	30.466833
<b>2</b>	55	50.013401	92	105.461264	30.193597
<b>3</b>	70	45.696322	92	113.461264	30.632114
<b>4</b>	53	50.504232	92	104.461264	29.889149

## Improving the model

```
In [62]: #Load the data  
cars_new = pd.read_csv("Cars.csv")
```

```
In [63]: #Discard the data points which are influencers and reassign the row number (reset_index)  
car1=cars_new.drop(cars_new.index[[70,76]],axis=0).reset_index()  
car1
```

```
Out[63]:
```

	index	HP	MPG	VOL	SP	WT
<b>0</b>	0	49	53.700681	89	104.185353	28.762059
<b>1</b>	1	55	50.013401	92	105.461264	30.466833
<b>2</b>	2	55	50.013401	92	105.461264	30.193597
<b>3</b>	3	70	45.696322	92	113.461264	30.632114
<b>4</b>	4	53	50.504232	92	104.461264	29.889149
<b>...</b>	...	...	...	...	...	...
<b>74</b>	75	175	18.762837	129	132.864163	42.778219
<b>75</b>	77	238	19.197888	115	150.576579	37.923113
<b>76</b>	78	263	34.000000	50	151.598513	15.769625
<b>77</b>	79	295	19.833733	119	167.944460	39.423099
<b>78</b>	80	236	12.101263	107	139.840817	34.948615

79 rows × 6 columns

```
In [33]: #Drop the original index  
car1=car1.drop(['index'],axis=1)
```

In [34]: `car1`

Out[34]:

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149
...	...	...	...	...	...
74	175	18.762837	129	132.864163	42.778219
75	238	19.197888	115	150.576579	37.923113
76	263	34.000000	50	151.598513	15.769625
77	295	19.833733	119	167.944460	39.423099
78	236	12.101263	107	139.840817	34.948615

79 rows × 5 columns

## Build Model

In [35]: `#Exclude variable "WT" and generate R-Squared and AIC values  
final_ml_V= smf.ols('MPG~VOL+SP+HP',data = car1).fit()  
final_ml_V.summary()`

Out[35]:

## OLS Regression Results

<b>Dep. Variable:</b>	MPG	<b>R-squared:</b>	0.816
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.809
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	111.0
<b>Date:</b>	Wed, 02 Feb 2022	<b>Prob (F-statistic):</b>	1.65e-27
<b>Time:</b>	20:29:10	<b>Log-Likelihood:</b>	-219.06
<b>No. Observations:</b>	79	<b>AIC:</b>	446.1
<b>Df Residuals:</b>	75	<b>BIC:</b>	455.6
<b>Df Model:</b>	3		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	25.5275	13.051	1.956	0.054	-0.471	51.526
<b>VOL</b>	-0.1825	0.023	-8.012	0.000	-0.228	-0.137
<b>SP</b>	0.4415	0.141	3.124	0.003	0.160	0.723
<b>HP</b>	-0.2291	0.035	-6.592	0.000	-0.298	-0.160
<b>Omnibus:</b>	6.541	<b>Durbin-Watson:</b>	1.130			
<b>Prob(Omnibus):</b>	0.038	<b>Jarque-Bera (JB):</b>	5.833			
<b>Skew:</b>	0.620	<b>Prob(JB):</b>	0.0541			
<b>Kurtosis:</b>	3.485	<b>Cond. No.</b>	5.76e+03			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.76e+03. This might indicate that there are strong multicollinearity or other numerical problems.


In [36]: `(final_ml_v.rsquared,final_ml_v.aic)`Out[36]: `(0.8161692010376005, 446.11722639447737)`In [37]: `#Exclude variable "VOL" and generate R-Squared and AIC values`  
`final_ml_w= smf.ols('MPG~WT+SP+HP',data = car1).fit()`In [38]: `(final_ml_w.rsquared,final_ml_w.aic)`Out[38]: `(0.8160034320495304, 446.18843235750313)`

Comparing above R-Square and AIC values, model 'final\_ml\_V' has high R- square and low AIC value hence include variable 'VOL' so that multi collinearity problem would be resolved.

## Cook's Distance

```
In [39]: model_influence_V = final_ml_V.get_influence()
(c_V, _) = model_influence_V.cooks_distance
```

```
In [40]: fig= plt.subplots(figsize=(20,7))
plt.stem(np.arange(len(car1)),np.round(c_V,3));
plt.xlabel('Row index')
plt.ylabel('Cooks Distance');
```



```
In [41]: #index of the data points where c is more than .5
(np.argmax(c_V),np.max(c_V))
```

```
Out[41]: (76, 1.1629387469135215)
```

```
In [42]: #Drop 76 and 77 observations
car2=car1.drop(car1.index[[76,77]],axis=0)
```

```
In [43]: car2
```

```
Out[43]:
```

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149
...	...	...	...	...	...
72	140	19.086341	160	124.715241	52.997752
73	140	19.086341	129	121.864163	42.618698
74	175	18.762837	129	132.864163	42.778219
75	238	19.197888	115	150.576579	37.923113
78	236	12.101263	107	139.840817	34.948615

77 rows × 5 columns

```
In [44]: #Reset the index and re arrange the row values
car3=car2.reset_index()
```

```
In [45]: car4=car3.drop(['index'],axis=1)
```

In [46]: `car4`

Out[46]:


	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149
...	...	...	...	...	...
72	140	19.086341	160	124.715241	52.997752
73	140	19.086341	129	121.864163	42.618698
74	175	18.762837	129	132.864163	42.778219
75	238	19.197888	115	150.576579	37.923113
76	236	12.101263	107	139.840817	34.948615

77 rows × 5 columns

In [47]: `#Build the model on the new data`  
`final_ml_V = smf.ols('MPG~VOL+SP+HP', data = car4).fit()`

In [48]: `#Again check for influencers`  
`model_influence_V = final_ml_V.get_influence()`  
`(c_V, _) = model_influence_V.cooks_distance`

In [49]: `fig= plt.subplots(figsize=(20,7))`  
`plt.stem(np.arange(len(car4)), np.round(c_V, 3));`  
`plt.xlabel('Row index')`  
`plt.ylabel('Cooks Distance');`



In [50]: `#index of the data points where c is more than .5`  
`(np.argmax(c_V), np.max(c_V))`

Out[50]: (65, 0.8774556986296811)

Since the value is  $<1$ , we can stop the diagnostic process and finalize the model

In [51]: `#Check the accuracy of the mode`



```
final_ml_V= smf.ols('MPG~VOL+SP+HP',data = car4).fit()
```

```
In [52]: final_ml_V.summary()
```

```
Out[52]:
```

#### OLS Regression Results

<b>Dep. Variable:</b>	MPG	<b>R-squared:</b>	0.867
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.861
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	158.6
<b>Date:</b>	Wed, 02 Feb 2022	<b>Prob (F-statistic):</b>	6.81e-32
<b>Time:</b>	20:29:10	<b>Log-Likelihood:</b>	-200.71
<b>No. Observations:</b>	77	<b>AIC:</b>	409.4
<b>Df Residuals:</b>	73	<b>BIC:</b>	418.8
<b>Df Model:</b>	3		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	25.2974	11.336	2.232	0.029	2.706	47.889
<b>VOL</b>	-0.1362	0.021	-6.366	0.000	-0.179	-0.094
<b>SP</b>	0.4335	0.122	3.560	0.001	0.191	0.676
<b>HP</b>	-0.2635	0.031	-8.634	0.000	-0.324	-0.203

<b>Omnibus:</b>	9.478	<b>Durbin-Watson:</b>	1.195
<b>Prob(Omnibus):</b>	0.009	<b>Jarque-Bera (JB):</b>	9.184
<b>Skew:</b>	0.770	<b>Prob(JB):</b>	0.0101
<b>Kurtosis:</b>	3.703	<b>Cond. No.</b>	5.72e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.72e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [53]: (final_ml_V.rsquared,final_ml_V.aic)
```

```
Out[53]: (0.8669636111859063, 409.4153062719508)
```

## Predicting for new data

```
In [54]: #New data for prediction
new_data=pd.DataFrame({'HP':40,"VOL":95,"SP":102},index=[1])
new_data
```

```
Out[54]:
```

	HP	VOL	SP
1	40	95	102

```
In [55]: final_ml_V.predict(new_data)
```

```
Out[55]: 1    46.035594  
dtype: float64
```

```
In [56]: final_ml_V.predict(cars_new.iloc[0:5,])
```

```
Out[56]: 0    45.428872  
1    43.992392  
2    43.992392  
3    43.508150  
4    44.085858  
dtype: float64
```

```
In [57]: cars_new.head()
```

```
Out[57]:
```

	HP	MPG	VOL	SP	WT
0	49	53.700681	89	104.185353	28.762059
1	55	50.013401	92	105.461264	30.466833
2	55	50.013401	92	105.461264	30.193597
3	70	45.696322	92	113.461264	30.632114
4	53	50.504232	92	104.461264	29.889149

```
In [58]: pred_y = final_ml_V.predict(cars_new)
```

```
In [59]: pred_y
```

```
Out[59]: 0    45.428872  
1    43.992392  
2    43.992392  
3    43.508150  
4    44.085858  
...  
76    7.165876  
77   12.198598  
78   14.908588  
79    4.163958  
80    9.161202  
Length: 81, dtype: float64
```

```
In [ ]:
```